

50.007 Machine Learning

Ensemble Models

Roy Ka-Wei Lee

Assistant Professor, DAI/ISTD, SUTD



SINGAPORE UNIVERSITY OF
TECHNOLOGY AND DESIGN

Outline

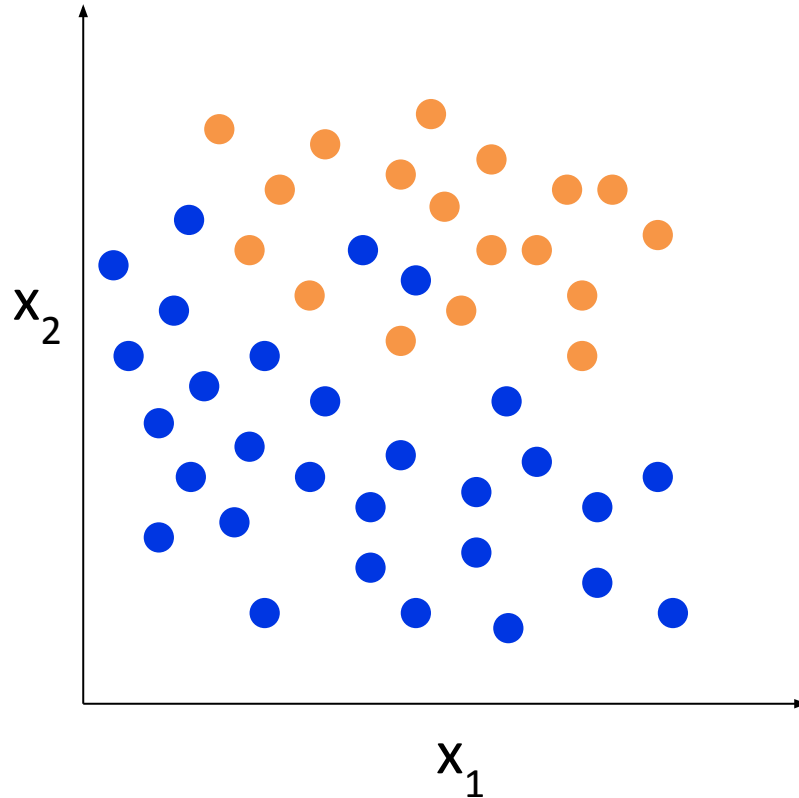
- Generalization
 - Underfitting and Overfitting
- Ensemble Classifiers
 - Bagging
 - Random Forest
 - Boosting
- Model Evaluation

Generalization

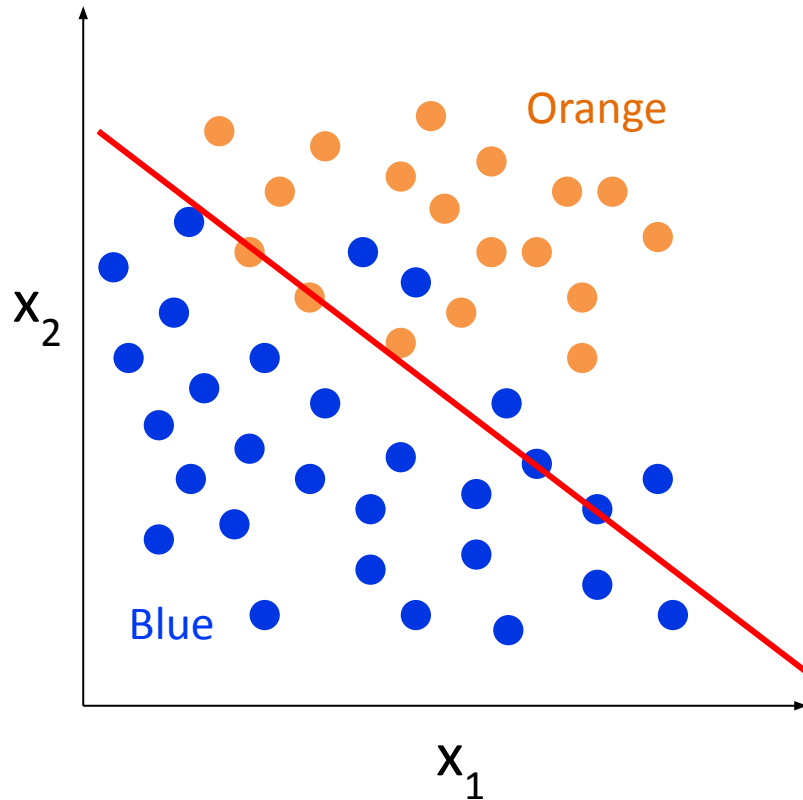
- Training data: $\{x_i, y_i\}$
 - Examples that we use to train our model
 - E.g. past records of days labelled with Roy goes/not go fishing
- Future/test data: $\{x_i, \hat{y}_i\}$
 - Examples that our model has not seen before
 - E.g. information about tomorrow to predict if Roy goes fishing
- Want to do well on future data not training data!
 - Not very useful to do well on training data; we already know the label
 - Easy to be perfect on training data
 - Doing well in training does not mean will do well on future data!

Classification Example

- Find a way to classify the points

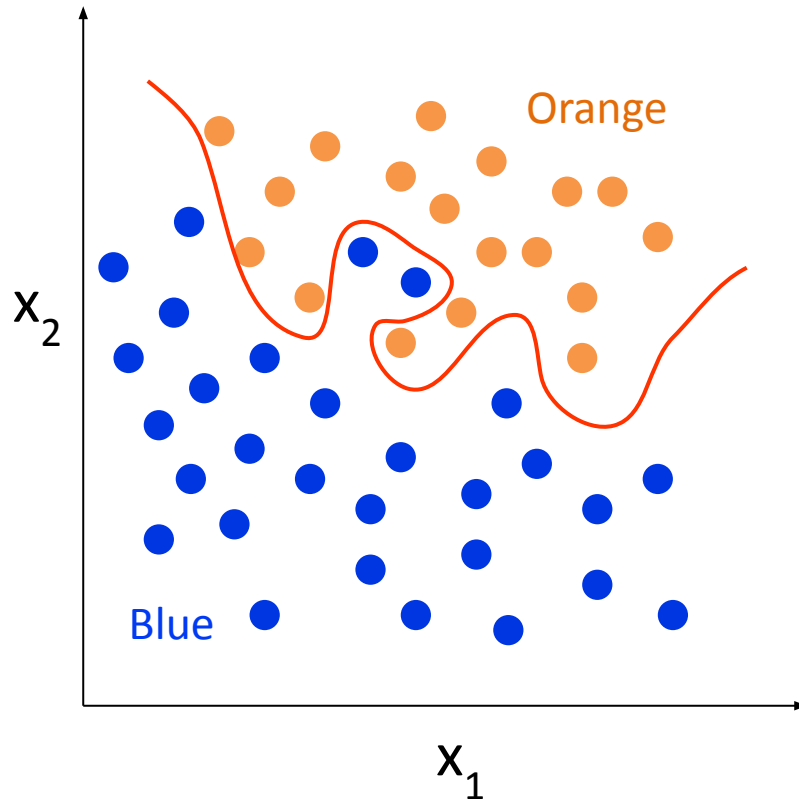


Classification Example



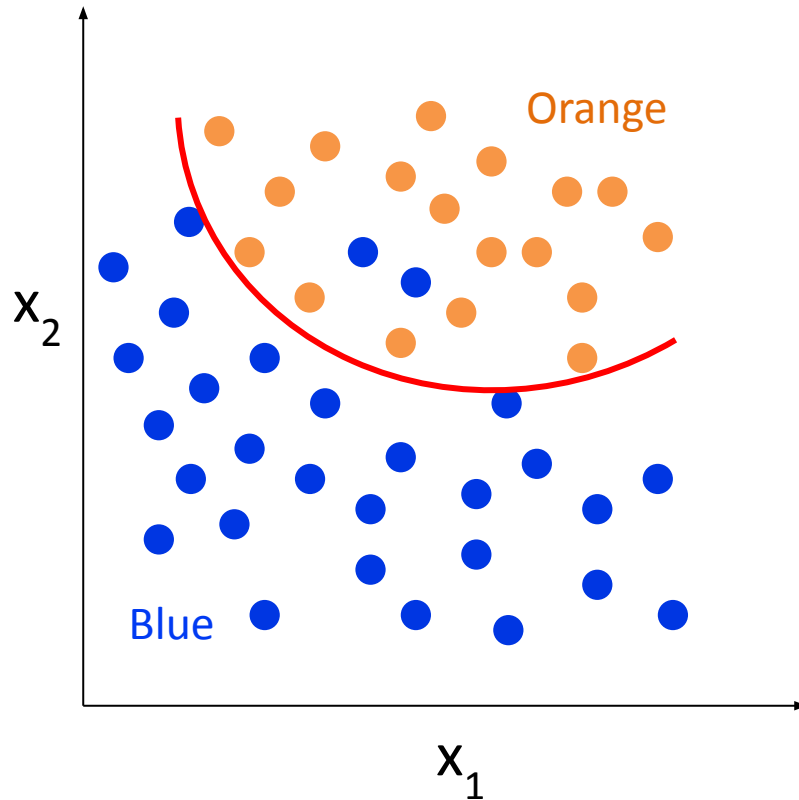
- Simple solution
- Unable to capture all salient pattern in data

Classification Example



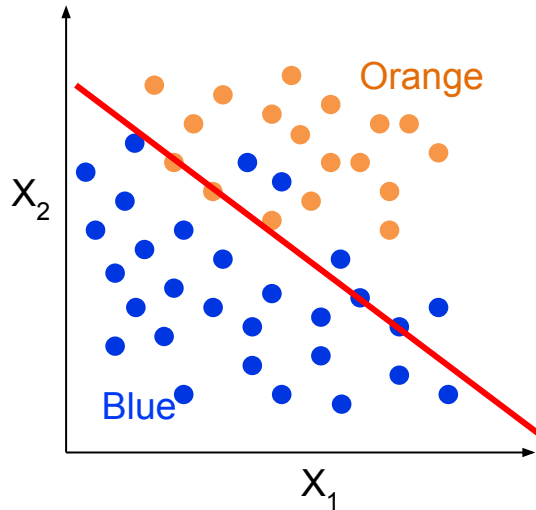
- Complex solution
- “Try too hard” to work well with training data
- Fit noise into the model
- Pattern is one-off and may not appear again in future data

Classification Example

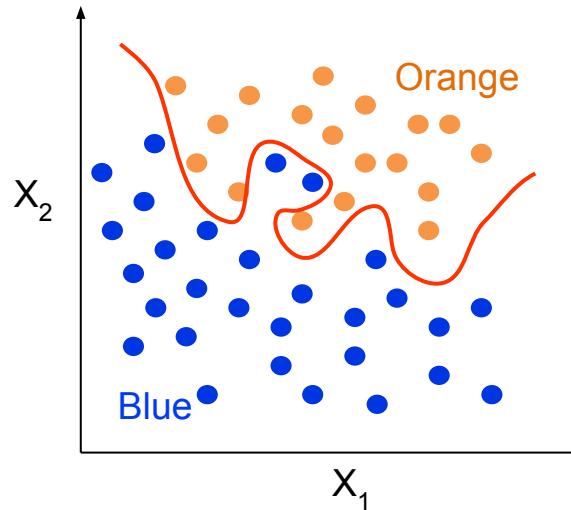


- Good solution, able to generalize well for future data

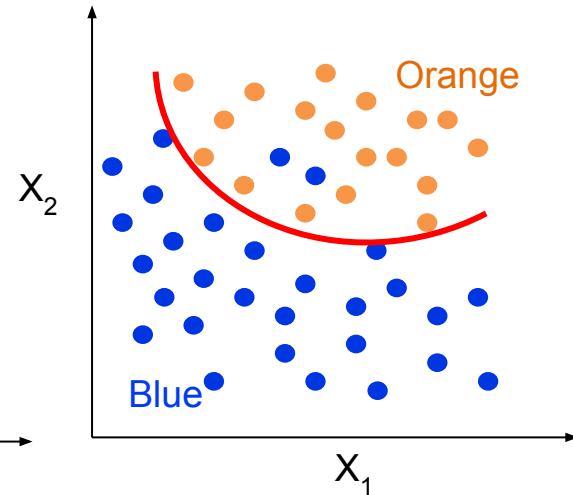
Underfitting and Overfitting



Underfitting
(Too simple to
explain the
variance)



Overfitting
(Too complex
- try too hard)

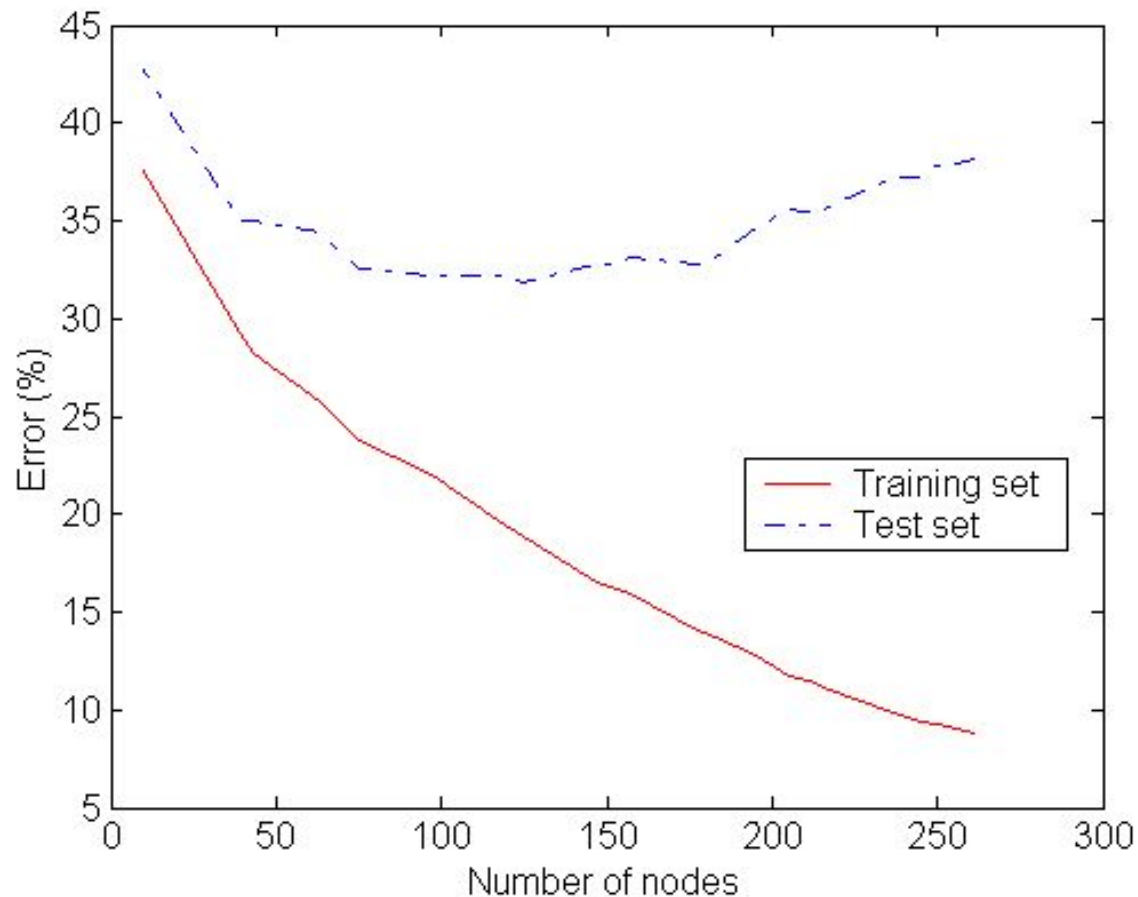


Appropriate
Fitting

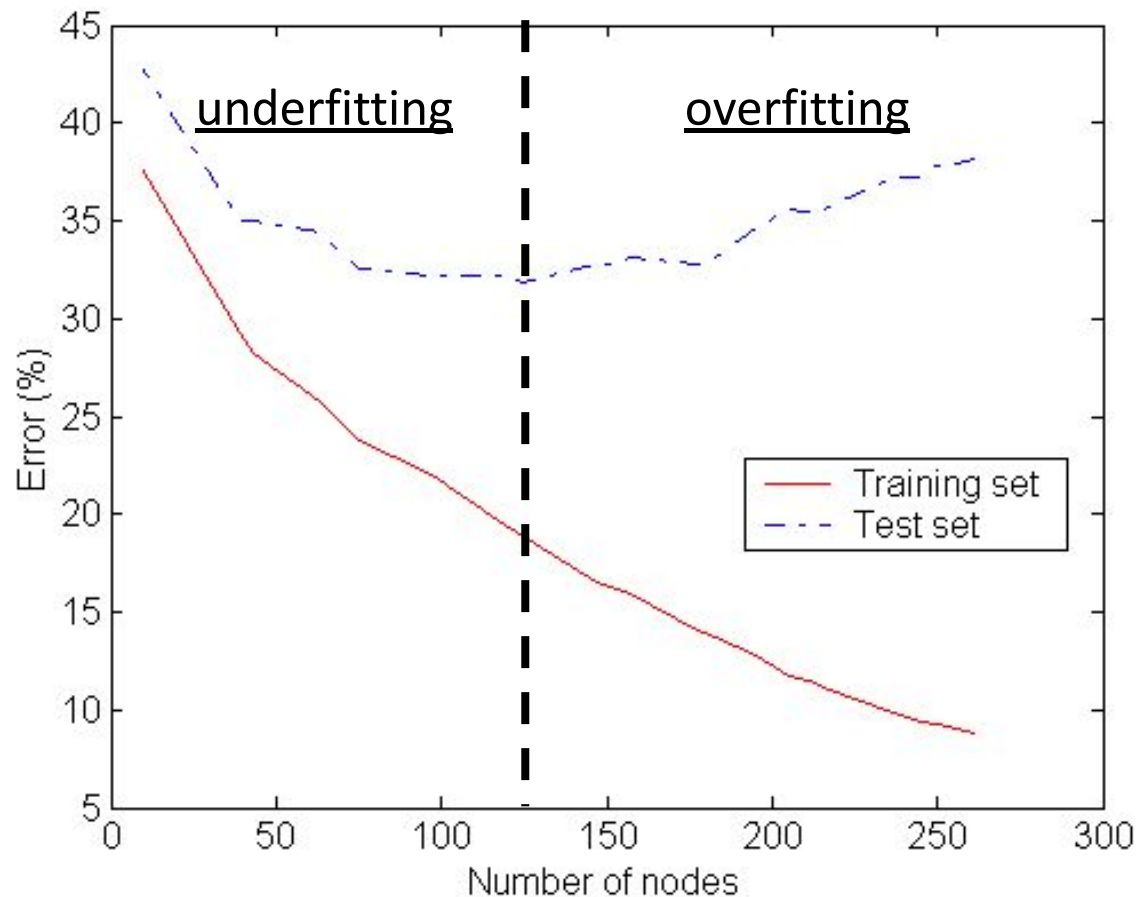
Underfitting and Overfitting

- Overfitting
 - Model is too complex (too flexible)
 - Fit noise in training data
 - Pattern is one-off and might not appear in future data
 - Model F overfits the data if:
 - We can find another model F'
 - Which makes more mistakes in training data: $E_{Train}(F') > E_{Train}(F)$
 - But less mistake in future data: $E_{Test}(F) > E_{Test}(F')$
- Underfitting
 - Model is too simple
 - Not powerful enough to capture the salient pattern
 - Can find another model F' with smaller $E_{Test}(F')$ and $E_{Train}(F')$

Underfitting and Overfitting

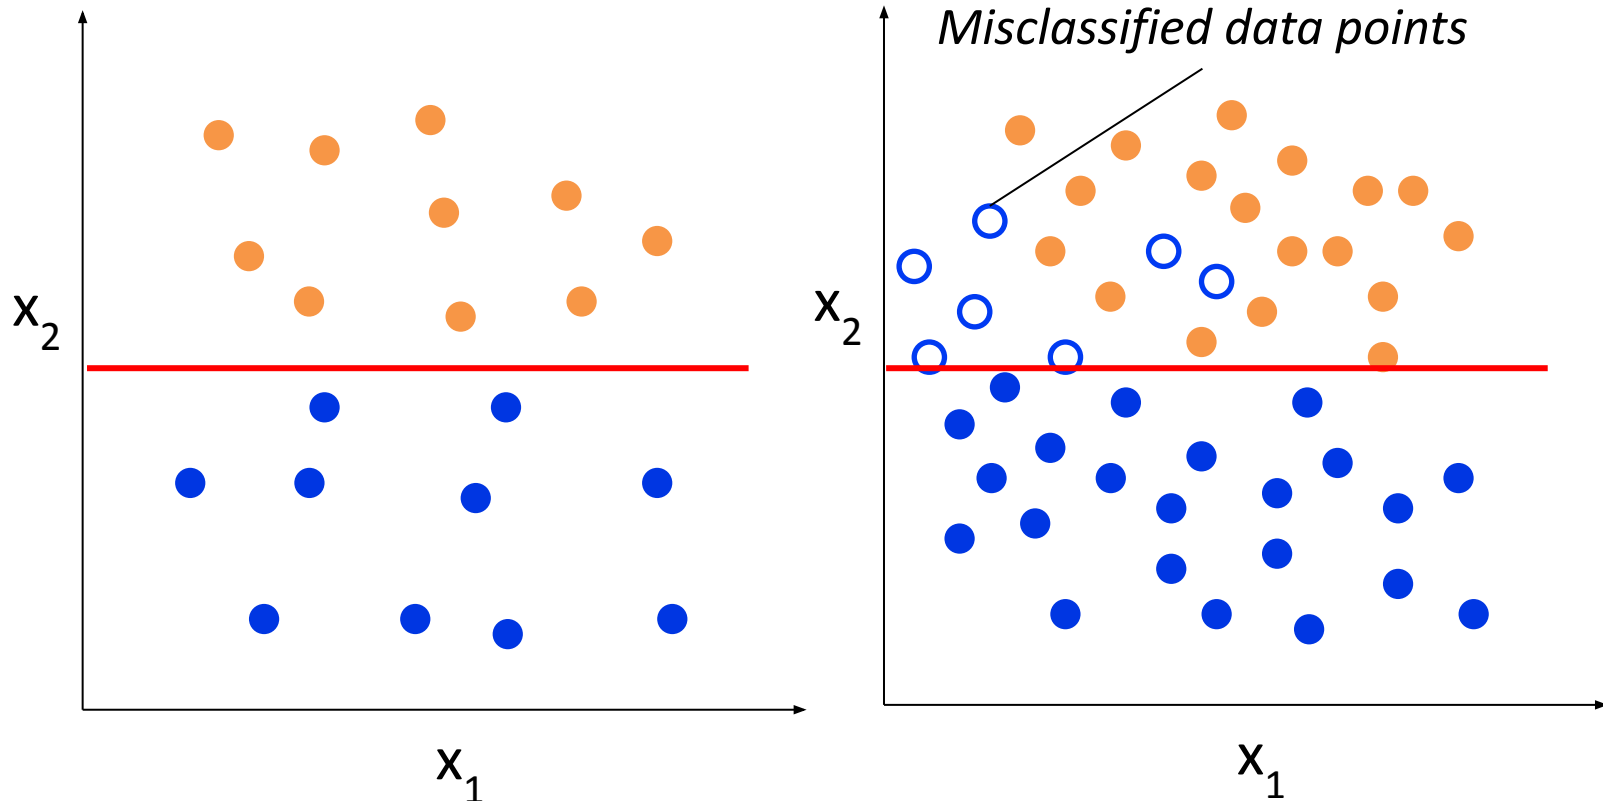


Underfitting and Overfitting



What Causes Overfitting?

- Noise in training data
- Insufficient data points in training data



Further Notes on Overfitting

- Overfitting results in decision trees that are more complex than necessary
- Training error no longer provides a good estimate of how well the tree will perform on previously unseen records
- Need new ways for estimating errors

Estimating Generalization Errors

- **Training errors:** error on training ($\sum e(t)$)
- **Generalization errors:** error on testing data ($\sum e'(t)$)
- Methods for **estimating** generalization errors:
 - **Optimistic approach:** simply use training error
 - $e'(t) = e(t)$
 - **Pessimistic approach:**
 - For each leaf node: $e'(t) = (e(t)+0.5)$
 - Total errors: $e'(T) = e(T) + N \times 0.5$ (N: number of leaf nodes)
 - Example: For a tree with 30 leaf nodes and 10 errors on training data (out of 1000 training data instances):
 - Training error = $10/1000 = 1\%$
 - Generalization error = $(10 + 30 \times 0.5)/1000 = 2.5\%$
 - **Reduced error pruning (REP):**
 - uses validation data set to estimate generalization error

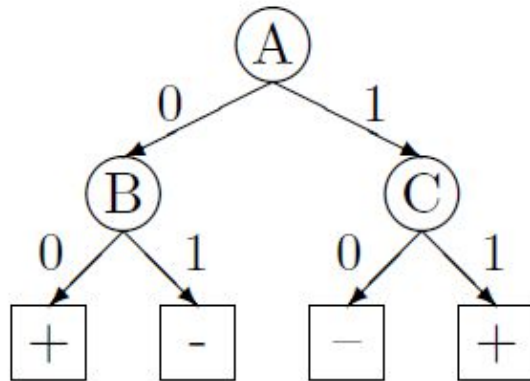
Estimating Generalization Errors

Training

S/N	A	B	C	Label
1	0	0	0	+
2	0	0	1	+
3	0	1	0	+
4	0	1	1	-
5	1	0	0	+

Testing

S/N	A	B	C	Label
6	0	0	0	+
7	0	1	1	+
8	1	1	0	+
9	1	0	1	-
10	1	0	0	+



Estimating Generalization Error

Optimistic: ?

Pessimistic: ?

Actual Generalization Error : ?

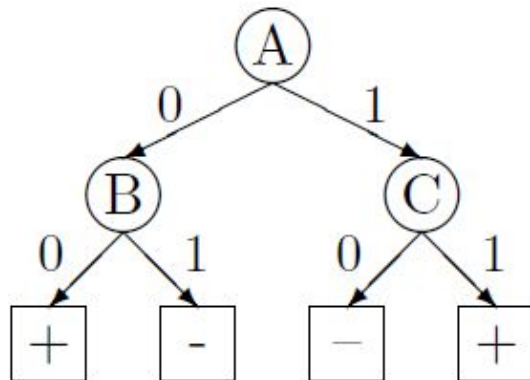
Estimating Generalization Errors

Training

S/N	A	B	C	Label	
1	0	0	0	+	✓
2	0	0	1	+	✓
3	0	1	0	+	✗
4	0	1	1	-	✓
5	1	0	0	+	✗

Testing

S/N	A	B	C	Label	
6	0	0	0	+	✓
7	0	1	1	+	✗
8	1	1	0	+	✗
9	1	0	1	-	✗
10	1	0	0	+	✗



Estimating Generalization Error

Optimistic: ?

Pessimistic: ?

Actual Generalization Error : ?

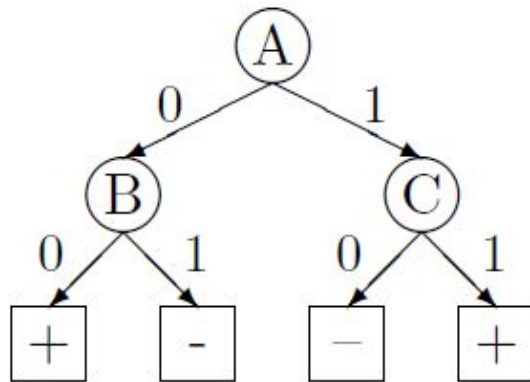
Estimating Generalization Errors

Training

S/N	A	B	C	Label	
1	0	0	0	+	✓
2	0	0	1	+	✓
3	0	1	0	+	✗
4	0	1	1	-	✓
5	1	0	0	+	✗

Testing

S/N	A	B	C	Label	
6	0	0	0	+	✓
7	0	1	1	+	✗
8	1	1	0	+	✗
9	1	0	1	-	✗
10	1	0	0	+	✗



Estimating Generalization Error

Optimistic: $e'(t) = e(t) = 2 / 5 = 0.4$

Pessimistic: $e'(t) = (e(t) + 0.5) = (2 + 4 * 0.5) / 5 = 0.8$

Actual Generalization Error : $e'(t) = 4 / 5 = 0.8$

Occam's Razor

- Given two models of similar generalization errors, one should prefer the simpler model over the more complex model
- For complex models, there is a greater chance that it was fitted accidentally by errors in data
- Therefore, one should include model complexity when evaluating a model

How to Address Overfitting

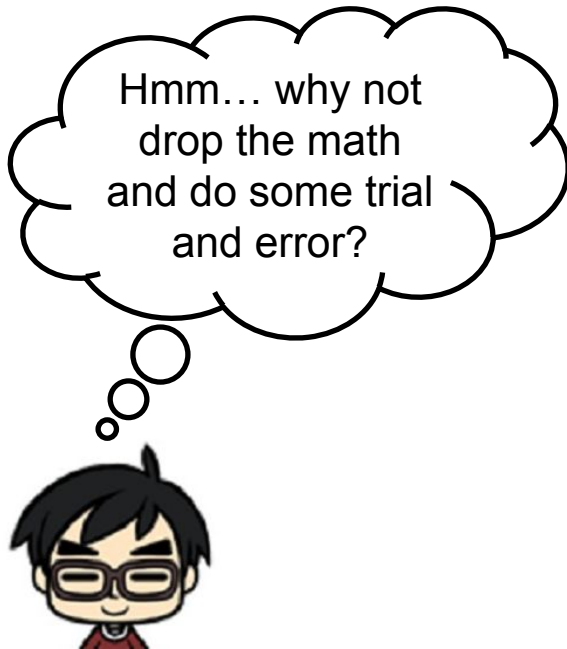
- **Pre-Pruning (Early Stopping Rule)**

- Stop the algorithm before it becomes a fully-grown tree
- Typical stopping conditions for a node:
 - Stop if all instances belong to the same class
 - Stop if all the attribute values are the same
 - More restrictive conditions: E.g., Stop if expanding the current node does not improve impurity measures (e.g., Gini or information gain).

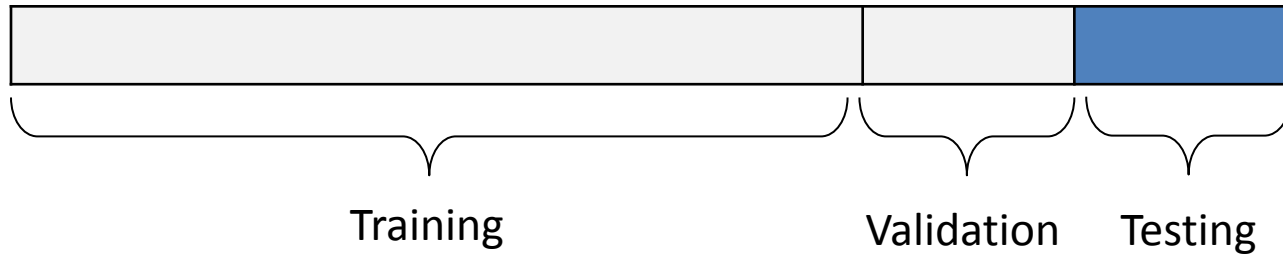
- **Post-Pruning**

- Grow decision tree to its entirety
- Trim the nodes of the decision tree in a bottom-up fashion
 - If generalization improves after trimming, replace sub-tree by a leaf node
- Class label of leaf node is determined from majority class of instances in the sub-tree

How will lazy Roy do it?



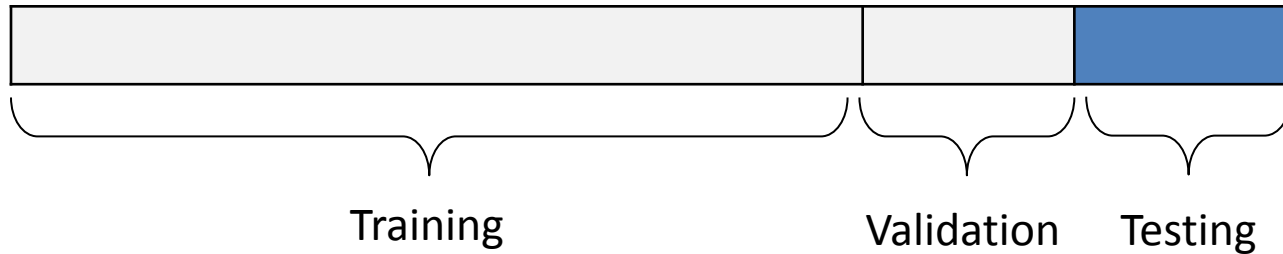
How will lazy Roy do it?



First split the
dataset into 3
parts...



How will lazy Roy do it?

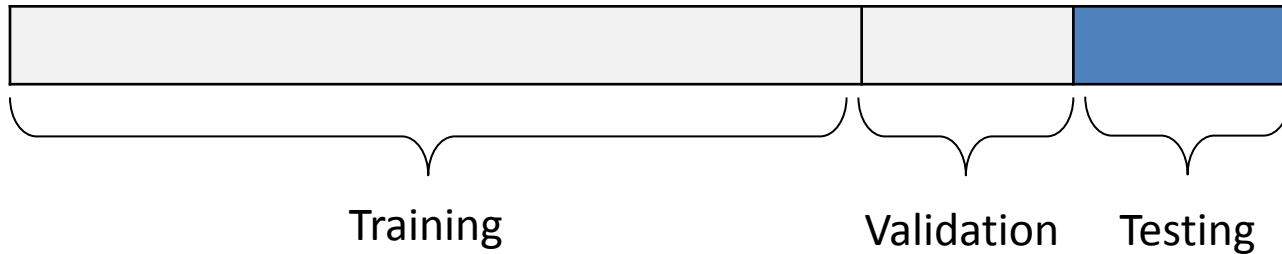


We train the model using the training data and test it on validation set...

```
DT = DecisionTreeClassifier(depth=d)  
DT.fit(training data)  
Pred = DT.predict(validation data)  
Acc = ComputeAcc(Pred)
```



How will lazy Roy do it?

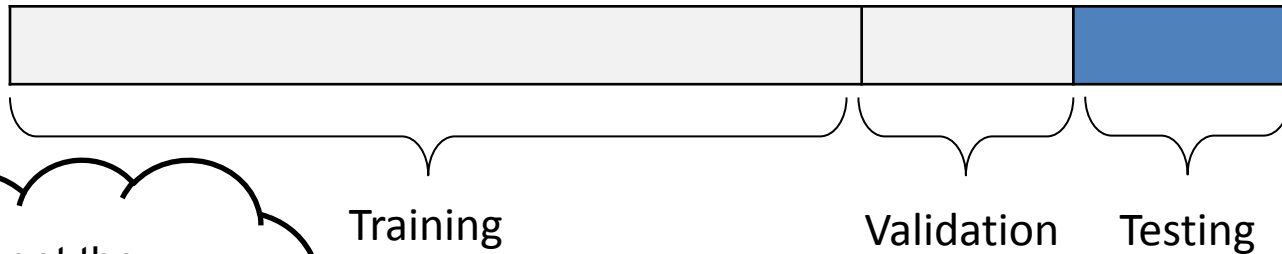


We loop and try
different depth and
get the best
accuracy

```
For d in range(101,1,-1)  
    DT = DecisionTreeClassifier(depth=d)  
    DT.fit(training data)  
    Pred = DT.predict(validation data)  
    Acc = ComputeAcc(Pred)  
    if BestAcc < Acc:  
        BestAcc = Acc  
        BestDepth = d
```



How will lazy Roy do it?



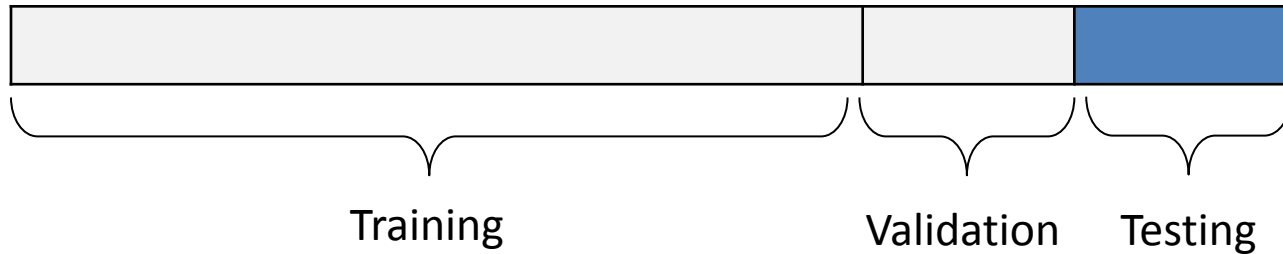
After we get the best depth, we re-train the model with it and test it on testing data

```
For d in range(101,1,-1)
    DT = DecisionTreeClassifier(depth=d)
    DT.fit(training data)
    Pred = DT.predict(validation data)
    Acc = ComputeAcc(Pred)
    if BestAcc < Acc:
        BestAcc = Acc
        BestDepth = d

DT = DecisionTreeClassifier(depth=BestDepth)
DT.fit(training data)
Pred = DT.predict(testing data)
Acc = ComputeAcc(Pred)
```



How will lazy Roy do it?



Caveat: workable solution but still have some drawbacks...

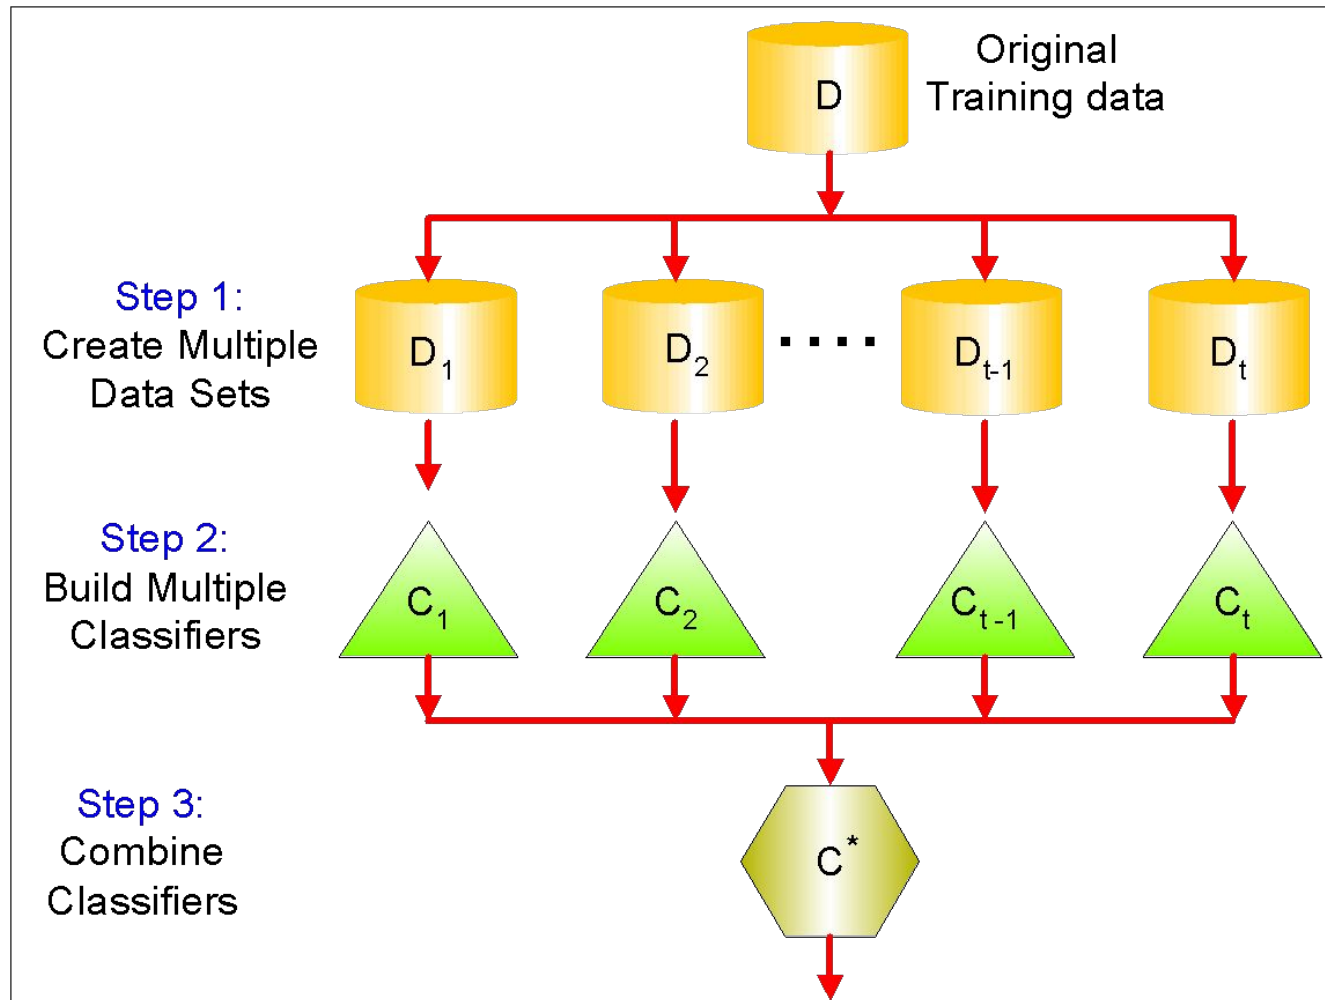


```
For d in range(101,1,-1)
    DT = DecisionTreeClassifier(depth=d)
    DT.fit(training data)
    Pred = DT.predict(validation data)
    Acc = ComputeAcc(Pred)
    if BestAcc < Acc:
        BestAcc = Acc
        BestDepth = d
DT = DecisionTreeClassifier(depth=BestDepth)
DT.fit(training data)
Pred = DT.predict(testing data)
Acc = ComputeAcc(Pred)
```

Ensemble Methods

- Construct a set of classifiers from the training data
- Predict class label of previously unseen records by aggregating predictions made by multiple classifiers
- Assumption:
 - Individual classifiers (voters) could be lousy (stupid), but the aggregate (electorate) can usually classify (decide) correctly.

General Idea



Why does it work?

- Suppose there are 25 base classifiers
 - Each classifier has error rate, $\varepsilon = 0.35$
 - Assume classifiers are independent
 - Probability that the ensemble classifier makes a wrong prediction (i.e., *13 out of the 25 classifiers misclassified*):

$$\sum_{i=13}^{25} \binom{25}{i} \varepsilon^i (1 - \varepsilon)^{25-i} = 0.06$$

Ensemble Methods

- How to generate an ensemble of classifiers?
(list below is non exhaustive!)
 - Bagging
 - Boosting

Bagging

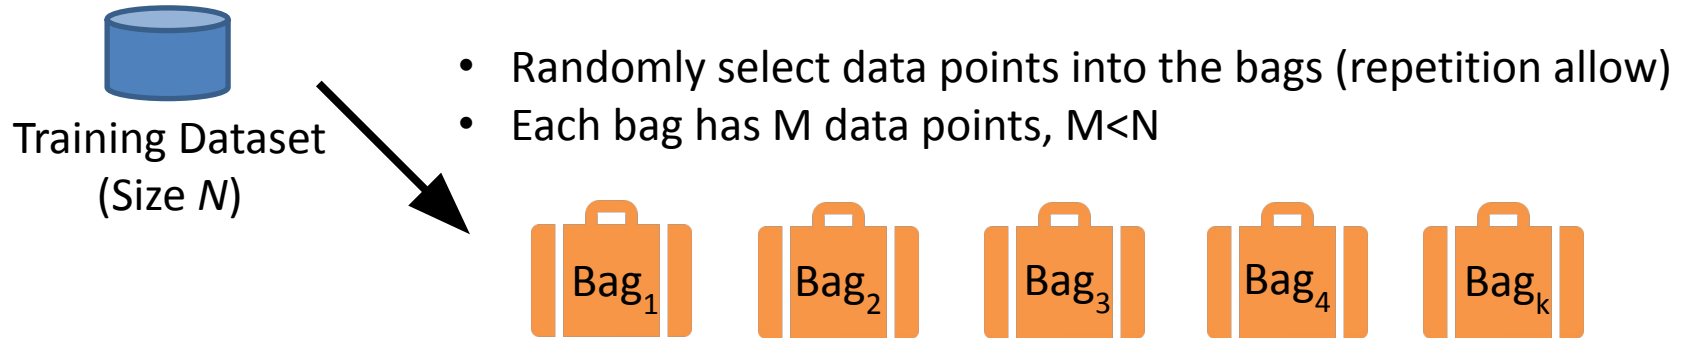
- **Bootstrap Aggregating (Bagging)**
- Multiple models of same learning algorithm trained with subset of dataset randomly sampled (with replacement) from the training dataset
- Each sample has probability $1 - (1 - 1/n)^n$ of being selected
 - This value tends to be **0.63** for large n

Bagging

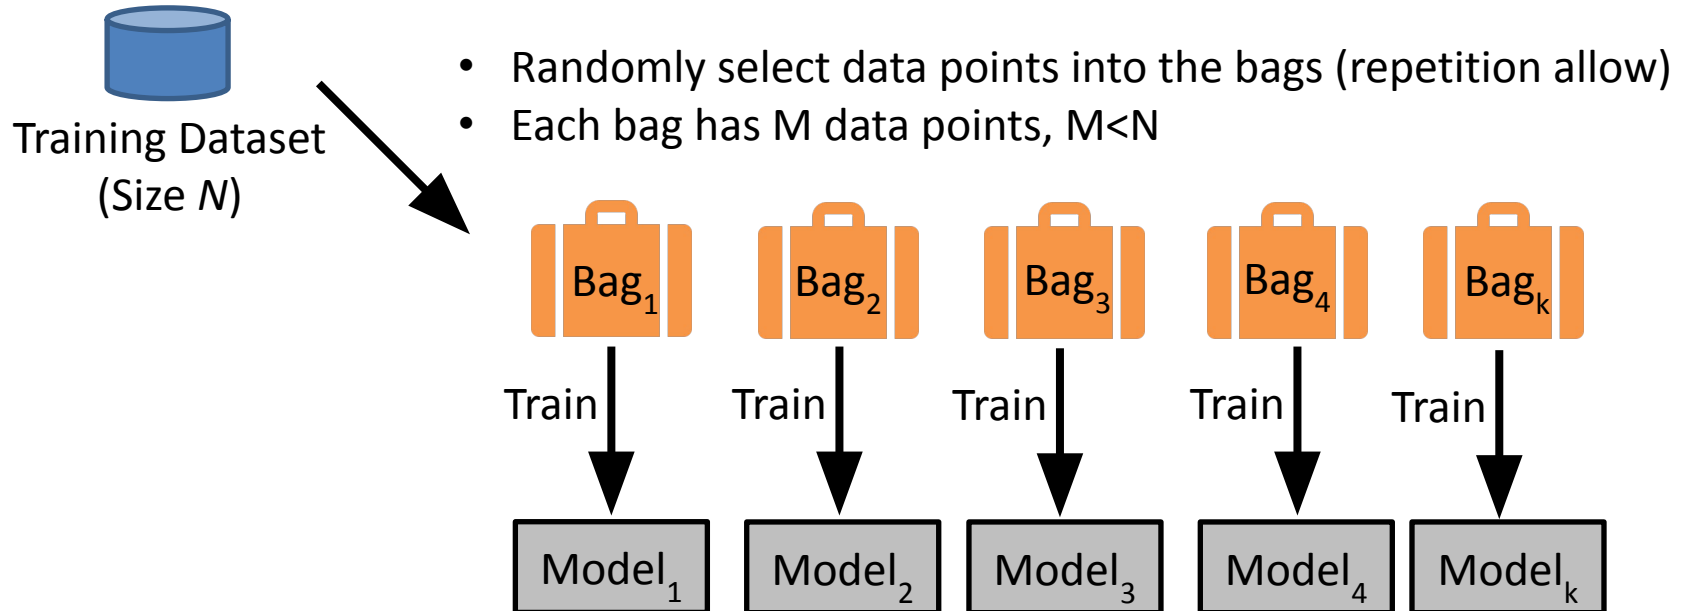


Training Dataset
(Size N)

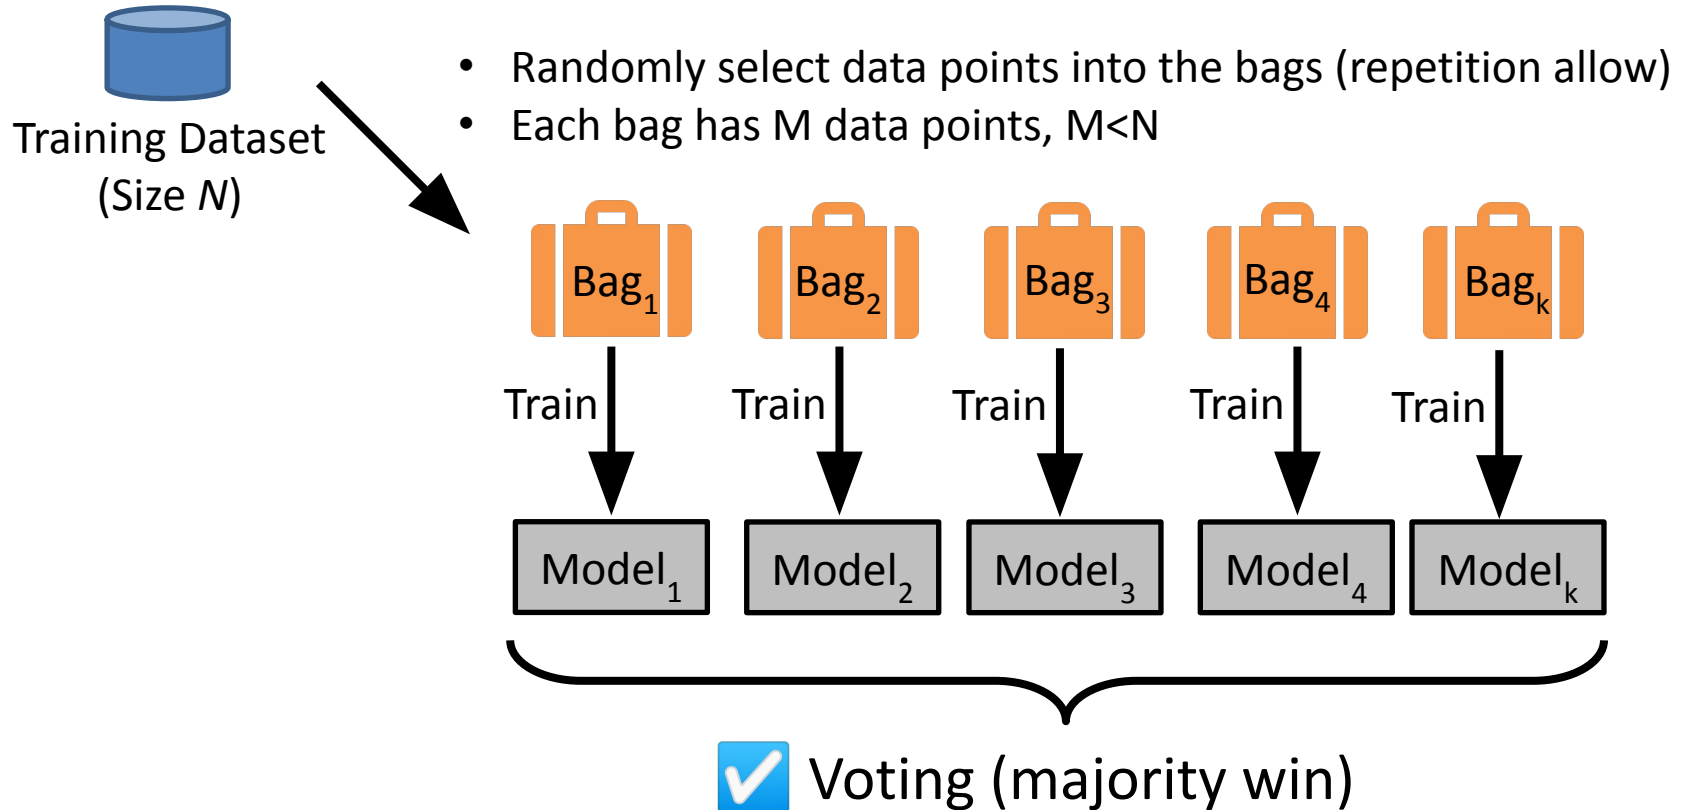
Bagging



Bagging



Bagging



Random Forest

- Train a lot of decision trees and use bagging
- Wisdom of the crowd!
 - Uncorrelated trees are preferred

Random Forest

- Train a lot of decision trees and use bagging
- Wisdom of the crowd!
 - Uncorrelated trees are preferred



Random Forest

- Bootstrapping + Feature Randomness



Boosting

- An iterative procedure to adaptively change distribution of training data by focusing more on previously misclassified records
 - Initially, all N records are assigned equal weights
 - Unlike bagging, sampling weights may change at the end of boosting round
 - Records that are **wrongly** classified will have their weights **increased**
 - Records that are **correctly** classified will have their weights **decreased**
 - **Caveat:** boosting show better predictive accuracy than bagging but also tends to over-fits the training data

Boosting

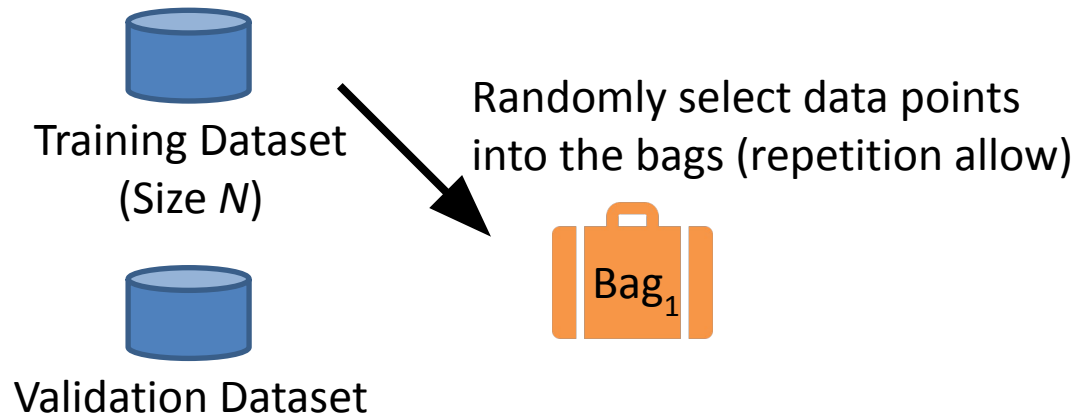


Training Dataset
(Size N)

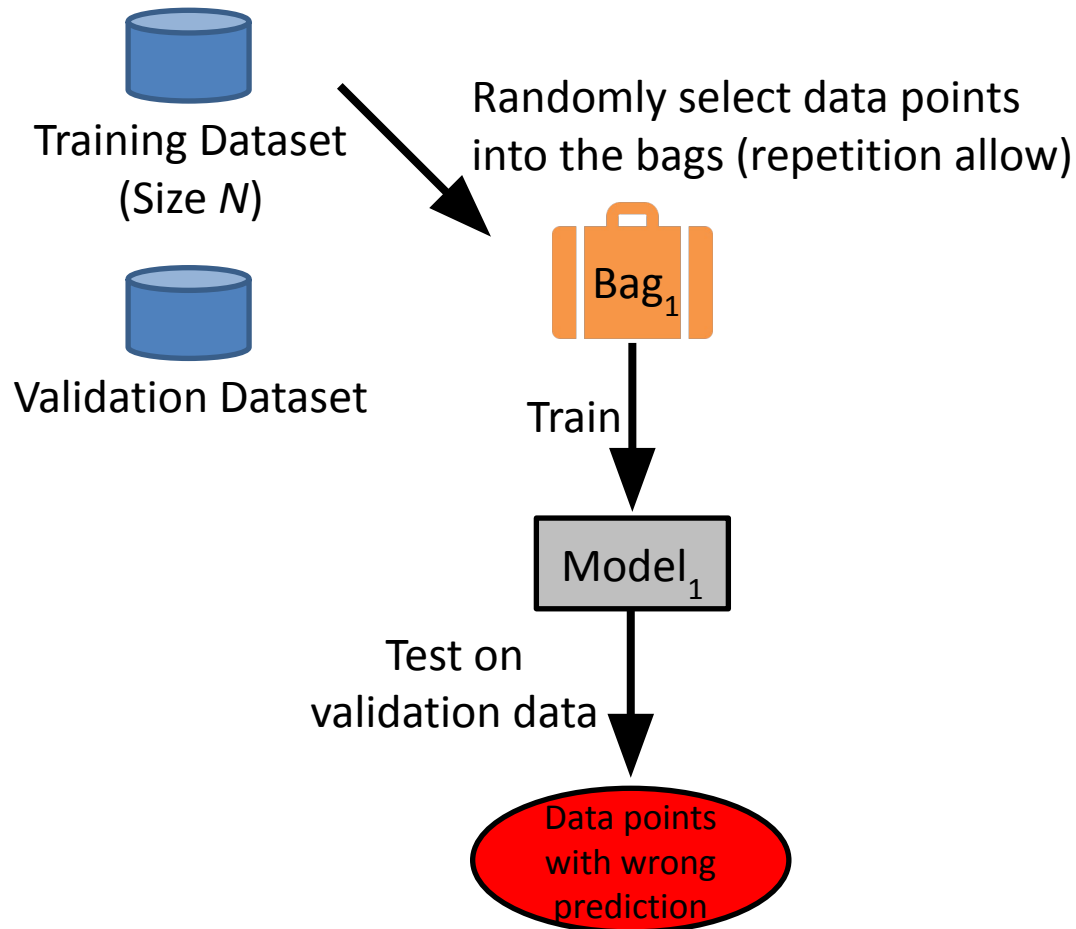


Validation Dataset

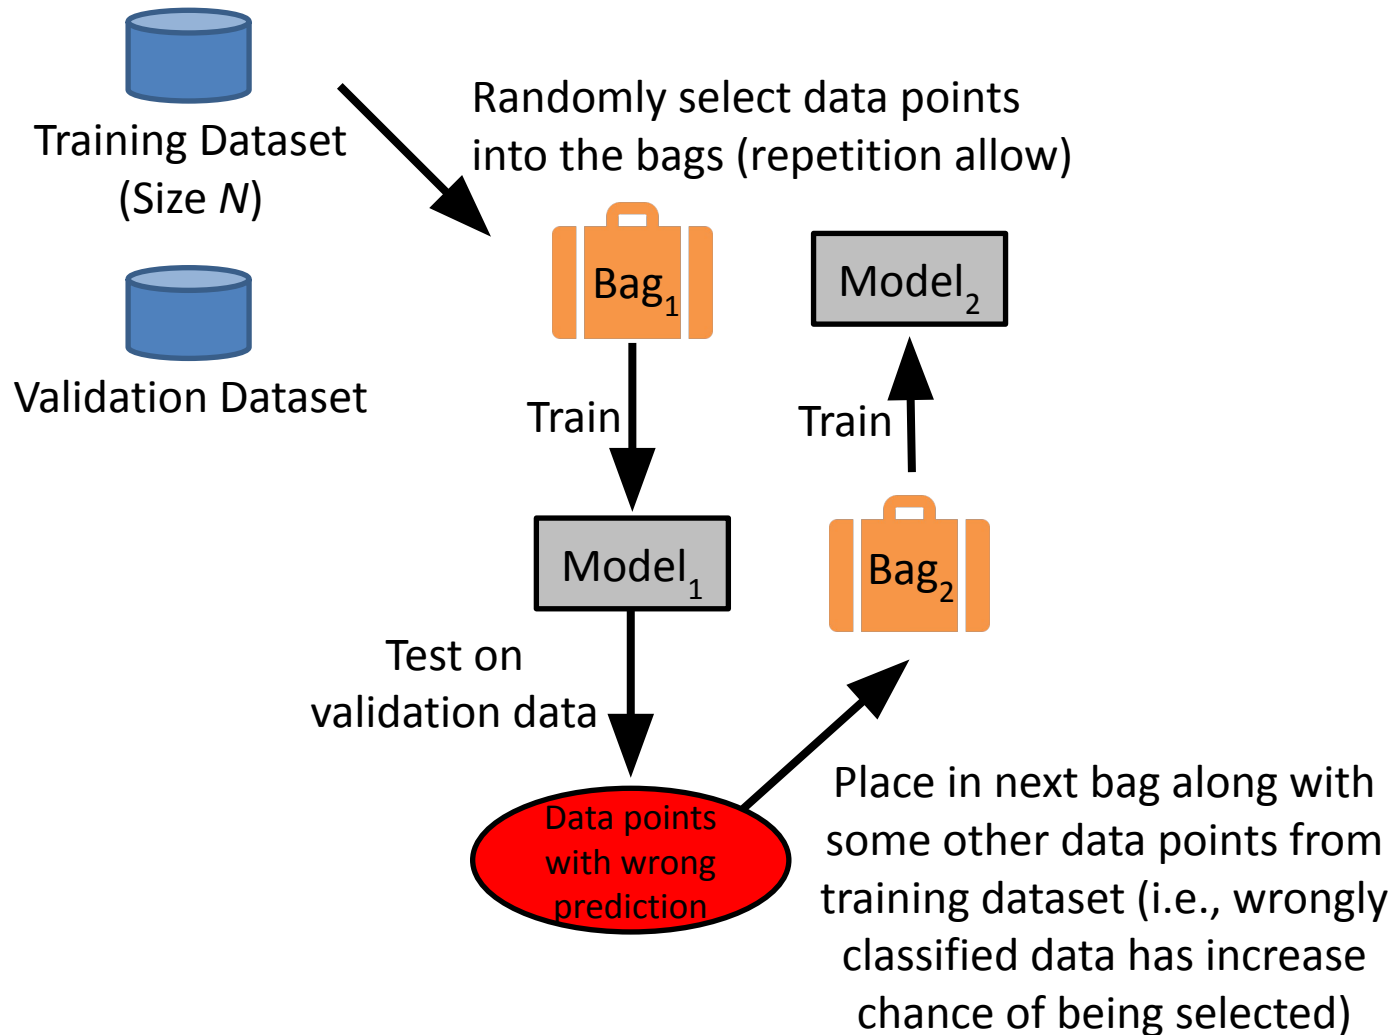
Boosting



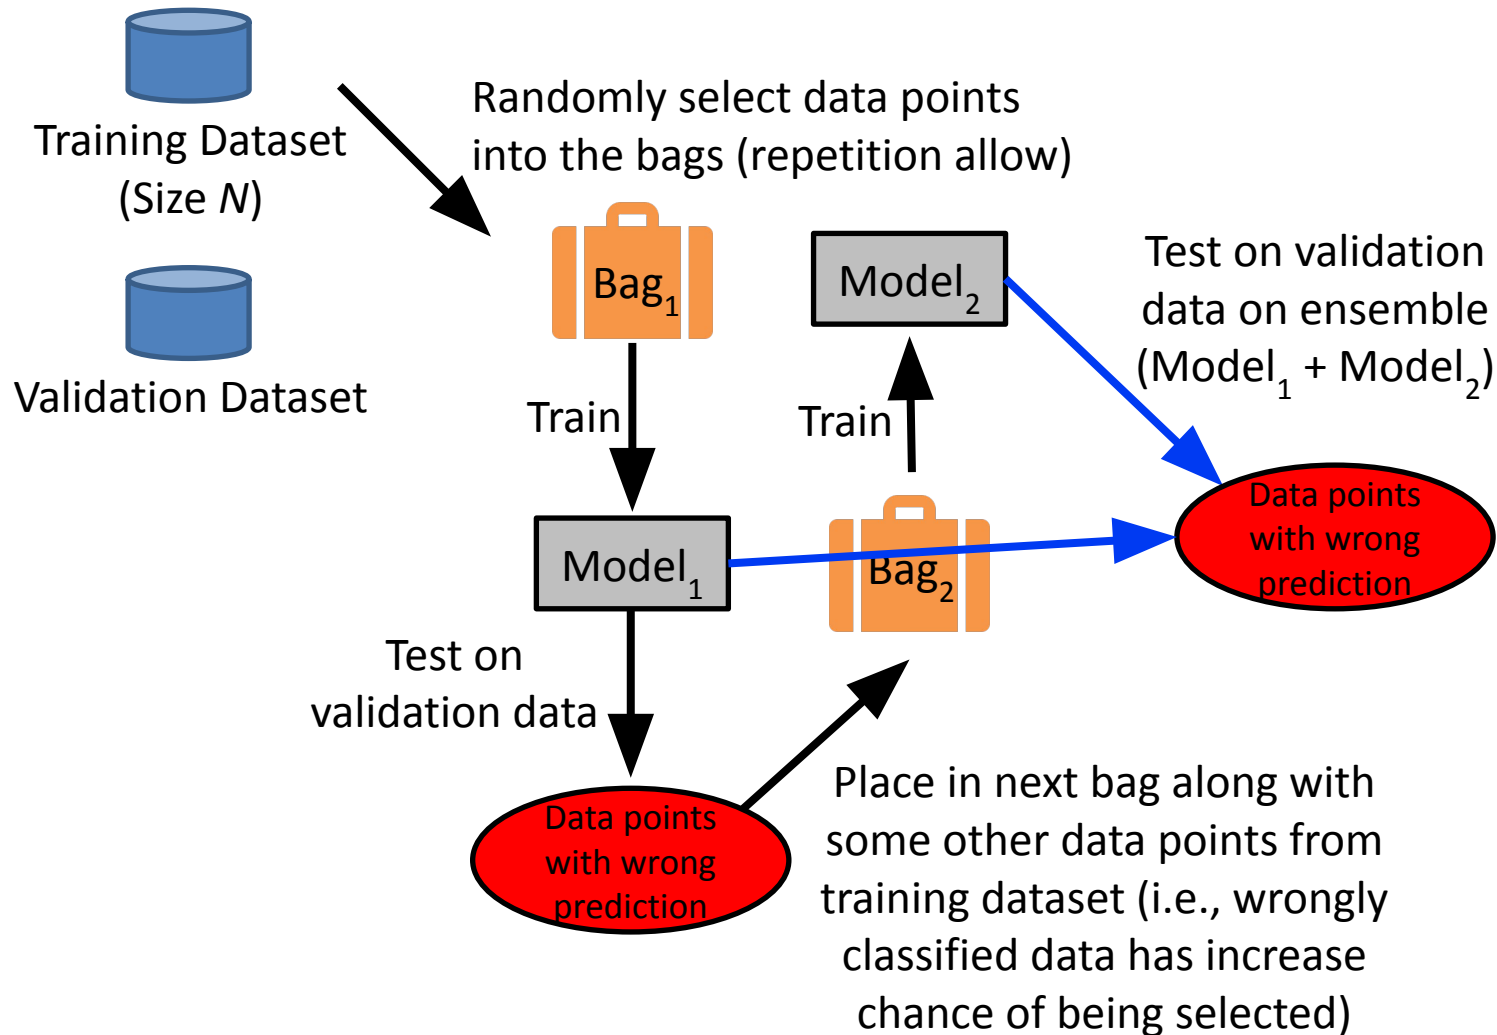
Boosting



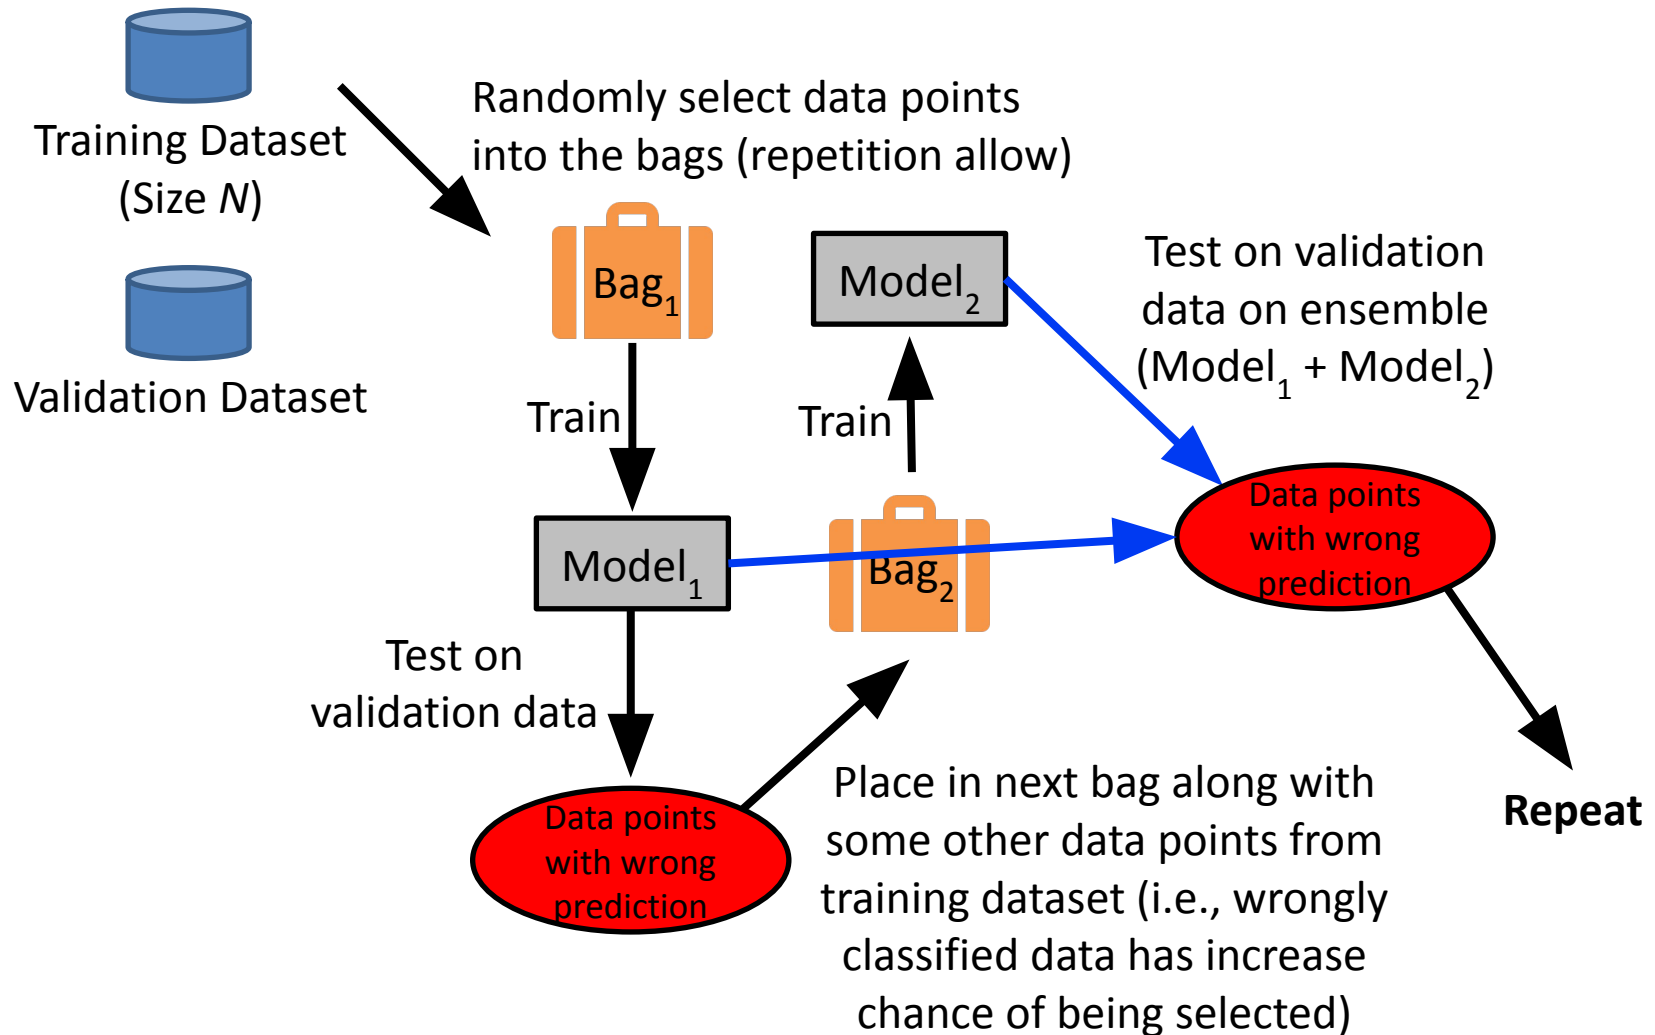
Boosting



Boosting

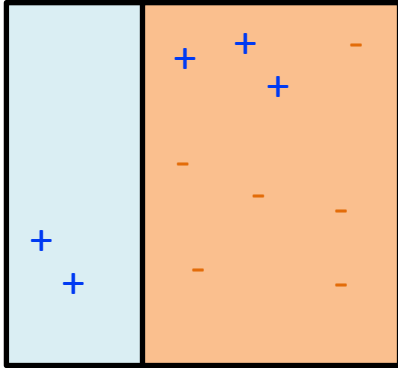


Boosting



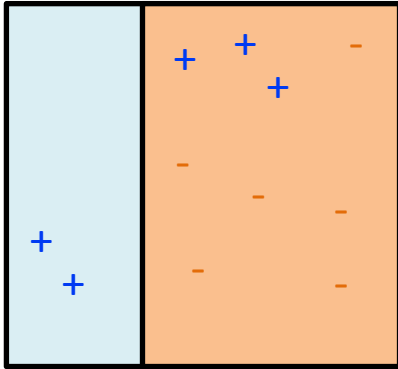
AdaBoost

Iteration 1

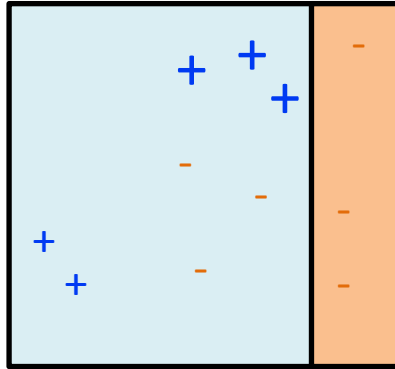


AdaBoost

Iteration 1

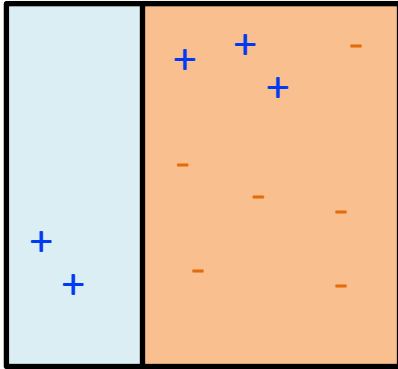


Iteration 2

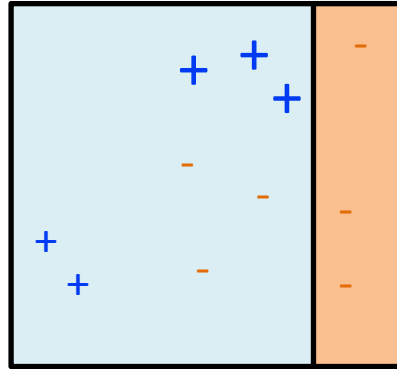


AdaBoost

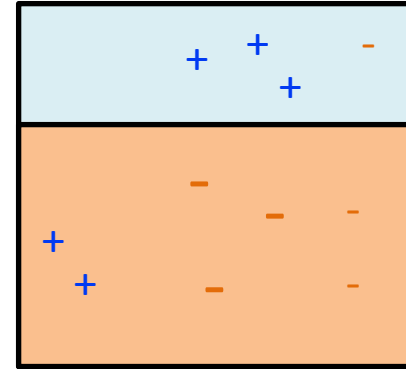
Iteration 1



Iteration 2

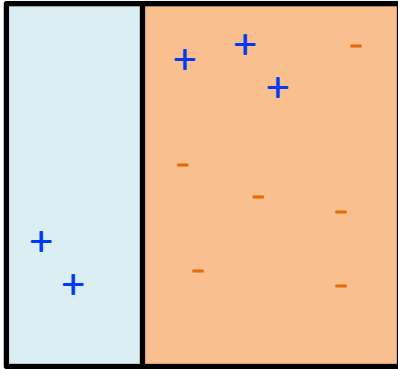


Iteration 3

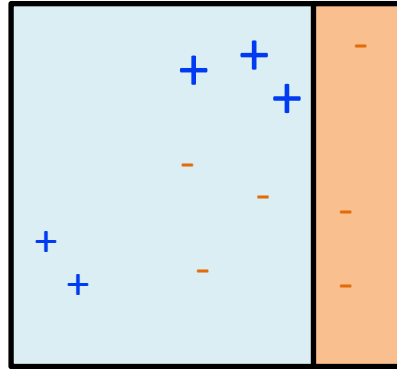


AdaBoost

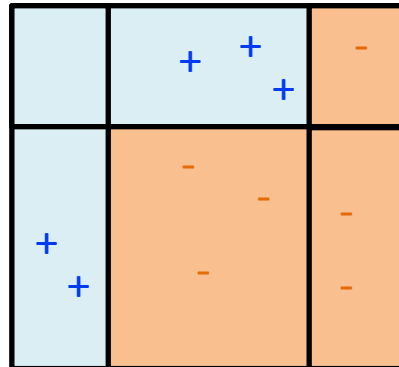
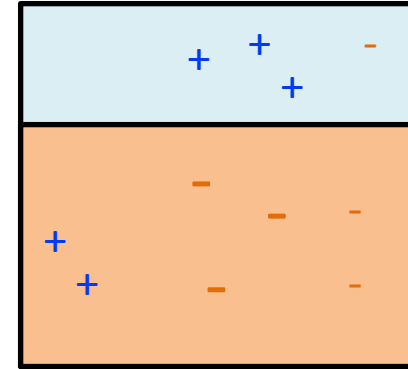
Iteration 1



Iteration 2



Iteration 3



Final Classifier/
Strong Classifier

AdaBoost Algorithm

- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: C_1, C_2, \dots, C_T

AdaBoost Algorithm

- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: C_1, C_2, \dots, C_T
- For $i=1$ to T :
 - Fit a classifier $C_i(x)$ to training data
 - Compute error rate of C_i :

$$\varepsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

AdaBoost Algorithm

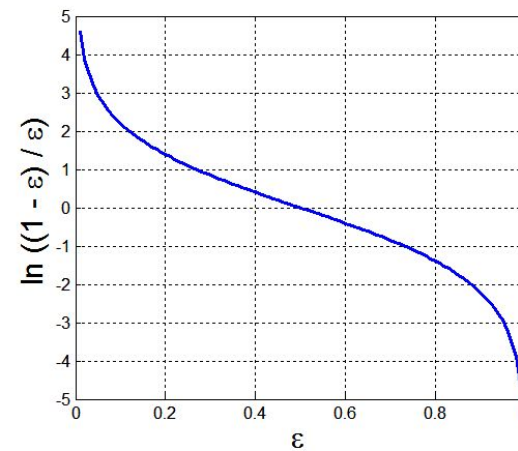
- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: C_1, C_2, \dots, C_T
- For $i=1$ to T :

- Fit a classifier $C_i(x)$ to training data
- Compute error rate of C_i :

$$\varepsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

- Compute importance of classifier C_i :

$$\alpha_i = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_i}{\varepsilon_i} \right)$$



When the error rate goes beyond 0.5, we assign the classifier a negative weight (i.e., we don't want them!)

AdaBoost Algorithm

- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: C_1, C_2, \dots, C_T
- For $i=1$ to T :
 - Fit a classifier $C_i(x)$ to training data
 - Compute error rate of C_i :

$$\varepsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

- Compute importance of classifier C_i :

$$\alpha_i = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_i}{\varepsilon_i} \right)$$

- Update the weights

$$w_i^{(j+1)} = \frac{w_i^{(j)}}{Z_j} \begin{cases} \exp^{-\alpha_j} & \text{if } C_j(x_i) = y_i \\ \exp^{\alpha_j} & \text{if } C_j(x_i) \neq y_i \end{cases}$$

where Z_j is the normalization factor

- $\exp^{-\alpha}$ is always lesser than 1 when the data point is correctly classified, thus giving it a smaller weight
- \exp^{α} is bigger than 1 when it is misclassified, assigning this data point a greater weight

AdaBoost Algorithm

- Initialize observation weights $w(x_i, y_i) = 1/n, i=1, \dots, n$
- Base classifiers: C_1, C_2, \dots, C_T
- For $i=1$ to T :

- Fit a classifier $C_i(x)$ to training data
- Compute error rate of C_i :

$$\varepsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

- Compute importance of classifier C_i :

$$\alpha_i = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_i}{\varepsilon_i} \right)$$

- Update the weights

$$w_i^{(j+1)} = \frac{w_i^{(j)}}{Z_j} \begin{cases} \exp^{-\alpha_j} & \text{if } C_j(x_i) = y_i \\ \exp^{\alpha_j} & \text{if } C_j(x_i) \neq y_i \end{cases}$$

where Z_j is the normalization factor

- Final Classifier: $C^*(x) = \arg \max_y \sum_{j=1}^T \alpha_j \delta(C_j(x) = y)$

Model Evaluation

- Metrics for Performance Evaluation
 - How to evaluate the performance of a model?
- Methods for Performance Evaluation
 - How to obtain reliable estimates?
- Methods for Model Comparison
 - How to compare the relative performance among competing models?

Metrics for Performance Evaluation

- Focus on the predictive capability of a model
 - Rather than how fast it takes to classify or build models, scalability, etc.
- Confusion Matrix:

	PREDICTED CLASS		
ACTUAL CLASS		Class=Yes	Class=No
	Class=Yes	a	b
	Class=No	c	d

a: TP (true positive)
b: FN (false negative)
c: FP (false positive)
d: TN (true negative)

Metrics for Performance Evaluation

ACTUAL CLASS	PREDICTED CLASS	
	Class=Yes	Class=No
	Class=Yes	Class=No
	a (TP)	b (FN)
	c (FP)	d (TN)

Most widely-used metric:

$$\text{Accuracy} = \frac{a + d}{a + b + c + d} = \frac{TP + TN}{TP + TN + FP + FN}$$

Limitation of Accuracy

- Consider a 2-class problem
 - Number of Class 0 examples = 9990
 - Number of Class 1 examples = 10
- If model predicts everything to be class 0, accuracy is $9990/10000 = 99.9\%$
 - Accuracy is misleading because model does not detect any class 1 example

Cost Matrix

- $C(i|j)$: Cost of misclassifying class j example as class i

	PREDICTED CLASS		
	$C(i j)$	Class=Yes	Class=No
ACTUAL CLASS	Class=Yes	C(Yes Yes) TP	C(No Yes) FN
	Class=No	C(Yes No) FP	C(No No) TN

TP: True Positive
TN: True Negative

FN: False Negative
FP: False Positive

Computing Cost of Classification

Cost Matrix	PREDICTED CLASS		
ACTUAL CLASS	C(i j)	+	-
	+	-1	100
	-	1	0

Model M_1	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	150	40
	-	60	250

Accuracy = ?

Cost = ?

Model M_2	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	250	45
	-	5	200

Accuracy = ?

Cost = ?

Computing Cost of Classification

Cost Matrix	PREDICTED CLASS		
ACTUAL CLASS	C(i j)	+	-
	+	-1	100
	-	1	0

Model M_1	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	150	40
	-	60	250

$$\text{Accuracy} = 150 + 250 / 150 + 250 + 40 + 60$$

$$= 80\%$$

$$\text{Cost} = 150(-1) + 40(100) + 60(1) + 250(0)$$

$$= 3910$$

Model M_2	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	250	45
	-	5	200

$$\text{Accuracy} = 250 + 200 / 250 + 200 + 45 + 5$$

$$= 90\%$$

$$\text{Cost} = 250(-1) + 45(100) + 5(1) + 200(0)$$

$$= 4255$$

Precision, Recall & F1

	PREDICTED CLASS		
ACTUAL CLASS		Class=Yes	Class=No
	Class=Yes	a (TP)	b (FN)
	Class=No	c (FP)	d (TN)

$$\text{Precision (p)} = \frac{a}{a + c} \quad \frac{\text{TP}}{\text{TP} + \text{FP}}$$

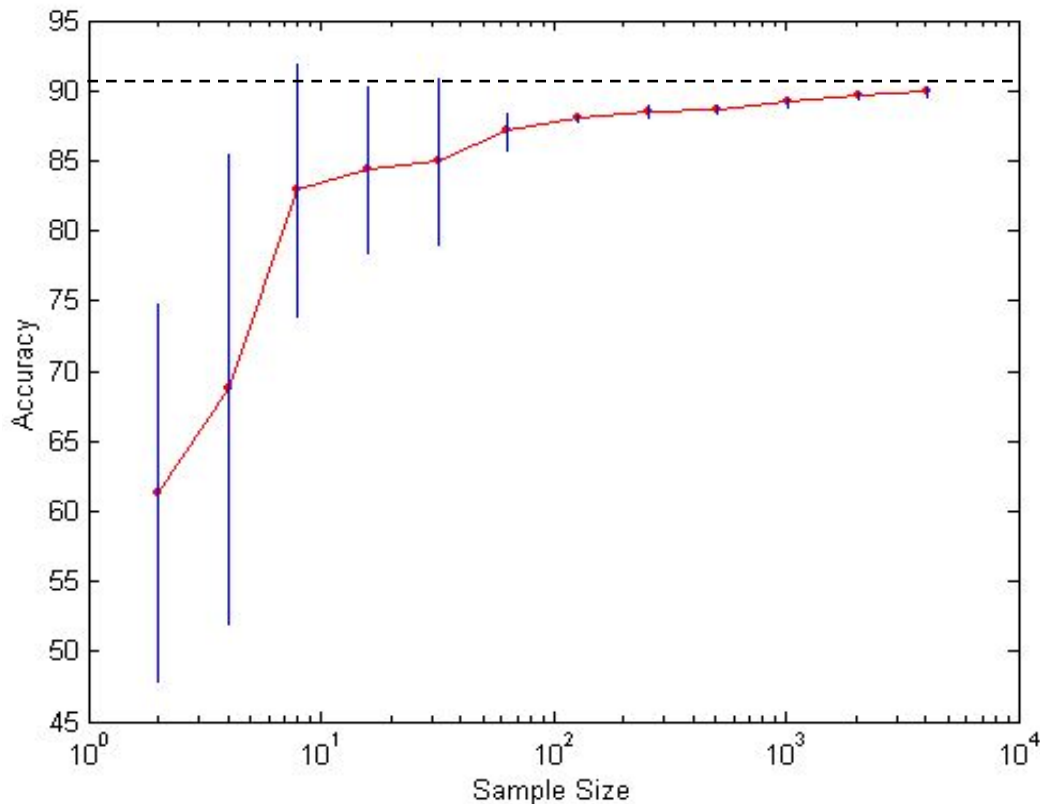
$$\text{Recall (r)} = \frac{a}{a + b} \quad \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{F - measure (F)} = \frac{2rp}{r + p} = \frac{2a}{2a + b + c}$$

Methods for Performance Evaluation

- How to obtain a reliable estimate of performance?
- Performance of a model may depend on other factors besides the learning algorithm:
 - Class distribution
 - Cost of misclassification
 - Size of training and test sets

Learning Curve



- Learning curve shows how accuracy changes with varying sample size
- Requires a sampling schedule for creating learning curve:
 - Arithmetic sampling (Langley, et al.)
 - Geometric sampling (Provost et al.)
- Effect of small sample size:
 - Bias in the estimate
 - Variance of estimate

Methods of Estimation

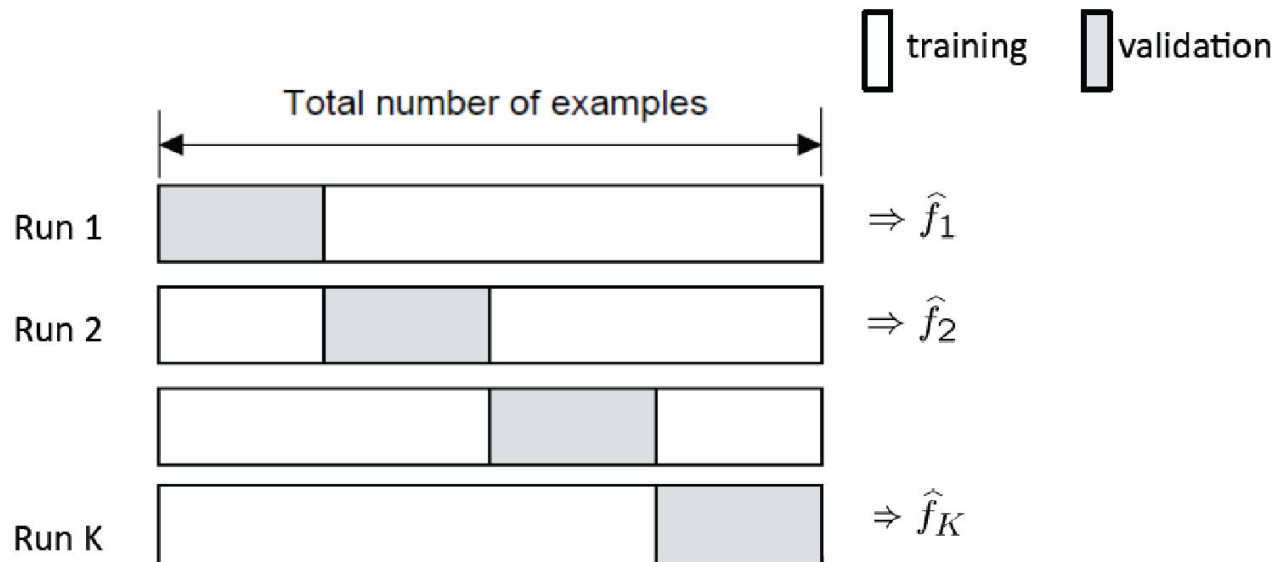
- Holdout
 - Reserve $2/3$ for training and $1/3$ for test
- Cross validation
 - Partition data into k disjoint subsets
 - k -fold: train on $k-1$ partitions, test on the remaining one
 - Leave-one-out: $k=n$
- Random subsampling
 - Repeated holdout

Train, Test and Validate

- **Training Set:** The actual dataset that we use to train the model (weights and biases in the case of Neural Network). The model sees and learns from this data.
- **Validation Set:** The sample of data used to provide an unbiased evaluation of a model fit on the training dataset while tuning model hyperparameters. The evaluation becomes more biased as skill on the validation dataset is incorporated into the model configuration.
- **Testing Set:** The sample of data used to provide an unbiased evaluation of a final model fit on the training dataset.

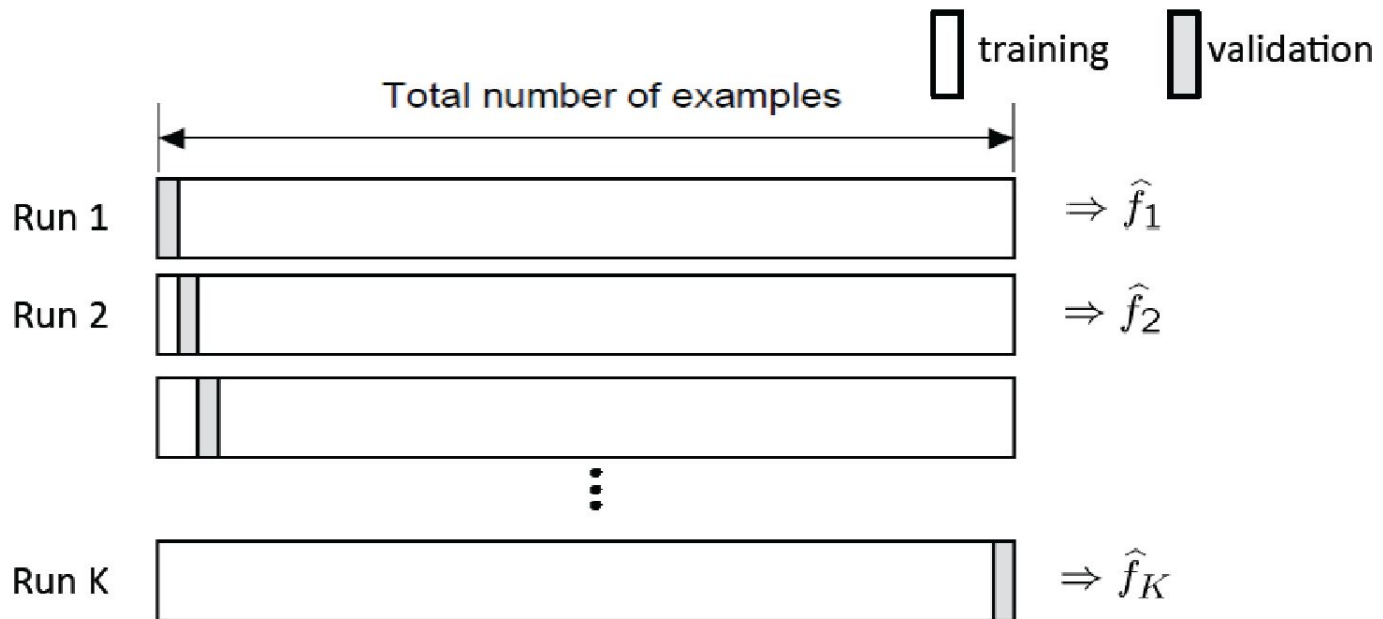
Cross Validation

- K-fold cross-validation
 - Create K-fold partition of the dataset.
 - Form K hold-out predictors, each time using one partition as validation and rest K-1 as training datasets.
 - Final predictor is average/majority vote over the K hold-out estimates.



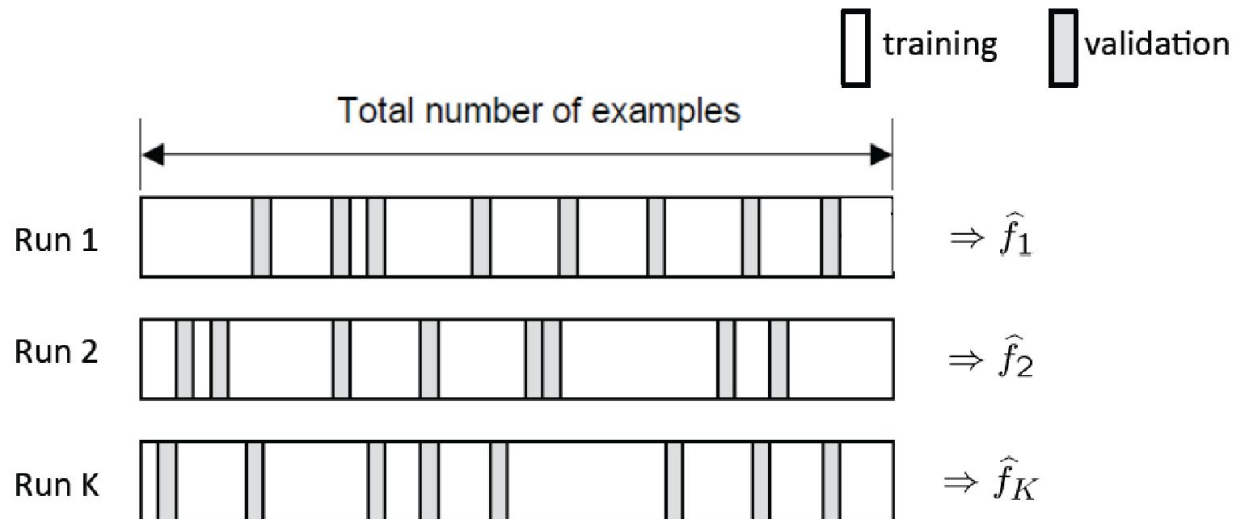
Cross Validation

- Leave-one-out (LOO) cross-validation
 - Special case of K-fold with $K=n$ partitions
 - Equivalently, train on $n-1$ samples and validate on only one sample per run for n runs



Random Subsampling

- Randomly subsample a fixed fraction α ($0 < \alpha < 1$) of the dataset for validation.
- Train the model with remaining data as training data.
- Repeat K times
- Compute the average metrics for the K hold-out estimates.



Summary

- Generalization
 - Underfitting and Overfitting
- Ensemble Classifiers
 - Bagging
 - Random Forest
 - Boosting
- Model Evaluation